

Hidden Topic Markov Models - Gruber et al

Alex Klibisz

September 29, 2016

Hidden Markov Models Overview

- ▶ System of tokens, assumed to be a Markov Process.
- ▶ State sequence is “hidden”, but tokens (which depend on state) are visible.
- ▶ Each state has a probability distribution over the possible tokens.
- ▶ Sequence of tokens generated by an HMM gives some information about the sequence of states.

Hidden Markov Models Example

Task: Given a sentence, determine the most likely sequence for its parts of speech.¹

- ▶ Some information is known based on prior data
 - ▶ States are parts of speech, tokens are the words.
 - ▶ $p(s'|s)$ - probability of transitioning from one state (part of speech) to another. Example: $p(\text{noun} \rightarrow \text{verb}) = 0.9$
 - ▶ $p(\text{word}|s)$ - probability of token (word) given a state. Example: $p(\text{"Blue"}|\text{noun}) = 0.4$
- ▶ Traverse the words to compute probability of each sequence.
- ▶ Sentence: *The blue bank closed.*
- ▶ $p(\text{det}, \text{adj}, \text{noun}, \text{verb}) =$
 $p(\text{"the"} \mid \text{det}) \cdot p(\text{adj} \mid \text{det}) \cdot p(\text{"blue"} \mid \text{adj}) \cdot \dots$

¹Source: <https://www.youtube.com/watch?v=7glSTzgjuwU>

Hidden Topic Markov Models Introduction

- ▶ Topics in a document are hidden and should be extracted.
- ▶ *Bag of words* is an unrealistic oversimplification.
- ▶ Topics should only transition at the beginning of a new sentence.
- ▶ Each document has a θ_d vector, representing its topic distribution.
- ▶ Topics transition based on binomial transition variable $\psi_n \in 0, 1$ for every word $w_1 \dots w_{N_d}$ in a document.

HTMM vs. LDA Visually

Abstract We give necessary and sufficient conditions for uniqueness of the **support** vector solution for the problems of pattern recognition and regression estimation, for a general class of cost functions. We show that if the solution is not unique, all **support** vectors are necessarily at bound, and we give some simple examples of non-unique solutions. We note that uniqueness of the primal (dual) solution does not necessarily imply uniqueness of the dual (primal) solution. We show how to compute the threshold b when the solution is unique, but when all **support** vectors are at bound, in which case the usual method for determining b does not work.

recognition and regression estimation algorithms [12], with arbitrary convex costs, the value of the normal w will always be unique. Acknowledgments C. Burges wishes to thank W. Keasler, V. Lawrence and C. Nohl of Lucent Technologies for their **support**. References [1] R. Fletcher. Practical Methods of Optimization. John Wiley and Sons, Inc., 2nd edition, 1987.

Abstract We give necessary and sufficient conditions for uniqueness of the **support** vector solution for the problems of pattern recognition and regression estimation, for a general class of cost functions. We show that if the solution is not unique, all **support** vectors are necessarily at bound, and we give some simple examples of non-unique solutions. We note that uniqueness of the primal (dual) solution does not necessarily imply uniqueness of the dual (primal) solution. We show how to compute the threshold b when the solution is unique, but when all **support** vectors are at bound, in which case the usual method for determining b does not work.

recognition and regression estimation algorithms [12], with arbitrary convex costs, the value of the normal w will always be unique. Acknowledgments C. Burges wishes to thank W. Keasler, V. Lawrence and C. Nohl of Lucent Technologies for their **support**. References [1] R. Fletcher. Practical Methods of Optimization. John Wiley and Sons, Inc., 2nd edition, 1987.

- HTMM (top) segments by sentence, LDA (bottom) segments only individual words.

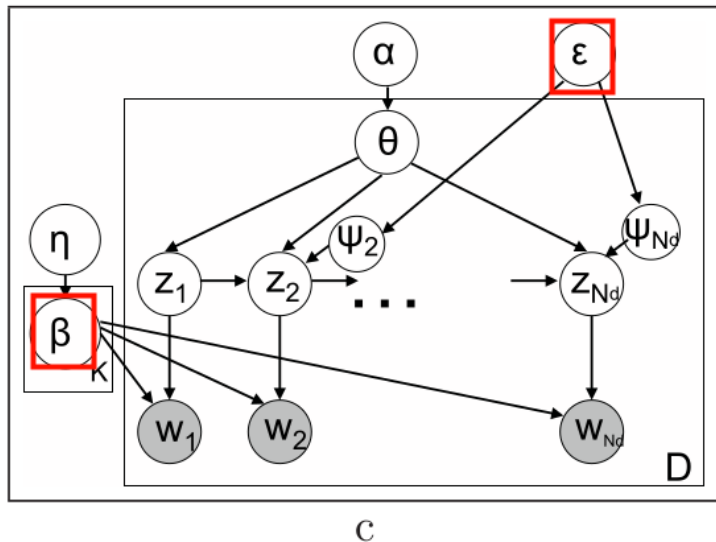
HTMM Definition

1. for $z=1\dots K$,
Draw $\beta_z \sim \text{Dirichlet}(\eta)$
2. for $d=1\dots D$,
Document d is generated as follows:
 - (a) Draw $\theta \sim \text{Dirichlet}(\alpha)$
 - (b) Set $\psi_1 = 1$
 - (c) for $n=2 \dots N_d$
 - i. If (begin_sentence) draw $\psi_n \sim \text{Binom}(\epsilon)$
else $\psi_n = 0$
 - (d) for $n=1 \dots N_d$
 - i. if $\psi_n == 0$ then $z_n = z_{n-1}$
else $z_n \sim \text{multinomial}(\theta)$
 - ii. Draw $w_n \sim \text{multinomial}(\beta_{z_n})$

HTMM Definition (Annotated)

- ▶ For every latent topic $z = 1 \dots K$, draw a $\beta_z \sim \text{Dirichlet}(\eta)$
- ▶ Generate each document $d = 1 \dots D$ as follows:
 - ▶ Draw a topic distribution $\theta \sim \text{Dirichlet}(\alpha)$
 - ▶ Word 1 is a *new* topic, $\psi_1 = 1$
 - ▶ For every word $n = 2 \dots N_d$:
 - ▶ If it's the first word in a sentence, draw $\psi_1 \sim \text{Binom}(\epsilon)$, otherwise no topic transition $\psi_1 = 0$
 - ▶ For every word $n = 1 \dots N_d$:
 - ▶ If $\psi_n = 0$, topic doesn't change: $z_n = z_{n-1}$.
 - ▶ Else draw new $z_n \sim \text{multinomial}(\theta)$
 - ▶ Draw $w_n \sim \text{multinomial}(\beta_{z_n})$

Parameter Approximation



Parameter Approximation (cont.)

Use Estimation-Maximization Algorithm (EM)

- ▶ EM for HMMs distinguishes between latent variables (topics z_n , transition variables ψ_n) and parameters.
- ▶ Estimation step uses Forward-Backward Algorithm

Unknown Parameters

- ▶ θ_d - topic distribution for each document
- ▶ β - used for multinomial word distributions
- ▶ ϵ - used for binomial topic transition variables

Known Parameters

- ▶ Based on prior research
- ▶ $\alpha = 1 + \frac{50}{K}$ - used for drawing $\theta \sim \text{Dirichlet}(\alpha)$
- ▶ $\eta = 1.01$ - used for drawing $\beta_z \sim \text{Dirichlet}(\eta)$

Experiment: NIPS Dataset

Data

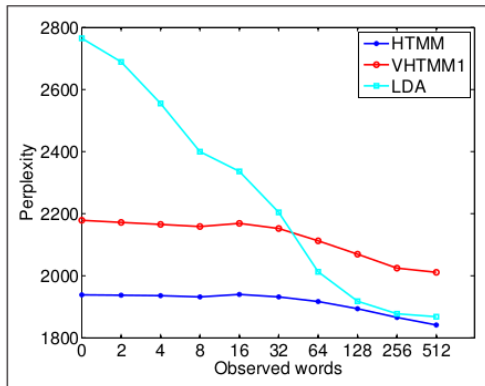
- ▶ 1740 documents, 1557 training, 183 testing.
- ▶ 12113 words in vocabulary.
- ▶ Extract vocabulary words, preserving order.
- ▶ Split sentences on punctuation . ? ! ;

Metric

- ▶ Perplexity for HTMM vs. LDA vs. VHTMM¹²
- ▶ *Perplexity reflects the difficulty of predicting a new unseen document after learning from a training set, lower is better.*

²VHTMM1 uses constant $\psi_n = 1$ so every sentence is of a new topic.

Figure 2: Perplexity as a function of observed words



- ▶ HTMM is significantly better than LDA for $N \leq 64$.
- ▶ Average document length is 1300 words.

Figure 3: topical segmentation in HTMM

Abstract We give necessary and sufficient conditions for uniqueness of the **support** vector solution for the problems of pattern recognition and regression estimation, for a general class of cost functions. We show that if the solution is not unique, all **support** vectors are necessarily at bound, and we give some simple examples of non-unique solutions. We note that uniqueness of the primal (dual) solution does not necessarily imply uniqueness of the dual (primal) solution. We show how to compute the threshold b when the solution is unique, but when all **support** vectors are at bound, in which case the usual method for determining b does not work.

recognition and regression estimation algorithms [12], with arbitrary convex costs, the value of the normal w will always be unique. Acknowledgments C. Burges wishes to thank W. Keasler, V. Lawrence and C. Nohl of Lucent Technologies for their **support**. References [1] R. Fletcher. Practical Methods of Optimization. John Wiley and Sons, Inc., 2nd edition, 1987.

- HTMM attributes “Support” to two different topics, mathematical and acknowledgments.


Figure 5: topical segmentation in LDA

Abstract We give necessary and sufficient conditions for uniqueness of the **support** vector solution for the problems of pattern recognition and regression estimation, for a general class of cost functions. We show that if the solution is not unique, all **support** vectors are necessarily at bound, and we give some simple examples of non-unique solutions. We note that uniqueness of the primal (dual) solution does not necessarily imply uniqueness of the dual (primal) solution. We show how to compute the threshold b when the solution is unique, but when all **support** vectors are at bound, in which case the usual method for determining b does not work.

recognition and regression estimation algorithms [12], with arbitrary convex costs, the value of the normal w will always be unique. **Acknowledgments** C. Burges wishes to thank W. Keasler, V. Lawrence and C. Nohl of Lucent Technologies for their **support**. **References** [1] R. Fletcher. *Practical Methods of Optimization*. John Wiley and Sons, Inc., 2nd edition, 1987.


- LDA attributes “Support” to only one topic.

Figure 6: Perplexity as a function of K



- Each topic limited to $N = 10$ words.

Figure 7: ϵ as a function of K



- ▶ Fewer topics \rightarrow lower $\epsilon \rightarrow$ infrequent transitions.
- ▶ More topics \rightarrow higher $\epsilon \rightarrow$ frequent transitions.

Table 1: Lower perplexity due to degrees of freedom?

- ▶ *Eliminate the option that the perplexity of HTMM might be lower than the perplexity of LDA only because it has less degrees of freedom (due to the dependencies between latent topics).*
- ▶ Generate and train on two datasets, $D = 1000$, $V = 200$, $K = 5$
- ▶ Dataset 1 generated with HTMM with $\epsilon = 0.1$ (likely to transition topics).
- ▶ Dataset 2 generated with LDA “bag of words”.
- ▶ HTMM still learns the correct parameters and outperforms on ordered data, does not outperform on “bag of words” data.
- ▶ Careful: *maybe bag of words was just a bad assumption for the NIPS dataset.*

Conclusion

Authors' Conclusion

- ▶ HTMM should be considered an extension of LDA.
- ▶ Markovian structure can learn more coherent topics, disambiguate topics of ambiguous words.
- ▶ Efficient learning and inference algorithms for HTMM already exist.

Questions

- ▶ Why only one dataset?
- ▶ How does it perform on unstructured speech? For example: run-on sentences or transcripts of live speech, debates, etc.
- ▶ Why is ψ_n necessary for every word, could you just consider the first words?

Thoughts for Project

Maybe shopping trips follow a Markov Model

- ▶ A shopper who purchases a crib on one store visit then diapers on the next visit might be having a baby soon.
- ▶ *These may be purchased with many other unrelated items, but there is still some meaning to them.*
- ▶ Extracting the crib, diapers, etc. sequence could help determine this meaning, whereas treating each shopping trip independently might ignore it.